

## QUES 2 : Kubernetes service types (ClusterIP, NodePort, LoadBalancer)

SOLN:

In Kubernetes, a **Service** is an abstraction that exposes a set of Pods as a network service. While Pods in Kubernetes are ephemeral and can change IPs, Services provide a stable endpoint for reliable communication.

Kubernetes supports **multiple types of Services** to handle different networking needs. These types determine **how and where your application can be accessed** — internally within the cluster or externally from the internet.

The 3 commonly used service types are:

### **ClusterIP (Default Service Type)**

- **Purpose:** Makes the service accessible only **within the cluster**.
- **Use Case:** Ideal for internal communication (e.g., frontend ↔ backend).
- **Example:** Backend database or internal microservices.

### **NodePort**

- **Purpose:** Exposes the service on a **static port** on each node's IP.
- **Use Case:** Allows external access by hitting <NodeIP>:<NodePort>.
- **Limitations:** Port range is limited (30000–32767), less secure for production.

### **LoadBalancer**

- **Purpose:** Provisions a **cloud load balancer** and assigns a public IP.
- **Use Case:** Best for production when deploying on cloud platforms like AKS, EKS, or GKE.
- **Benefit:** External users can access your service through a single public IP.

## My Deployment Using LoadBalancer (As Part of This Task) :

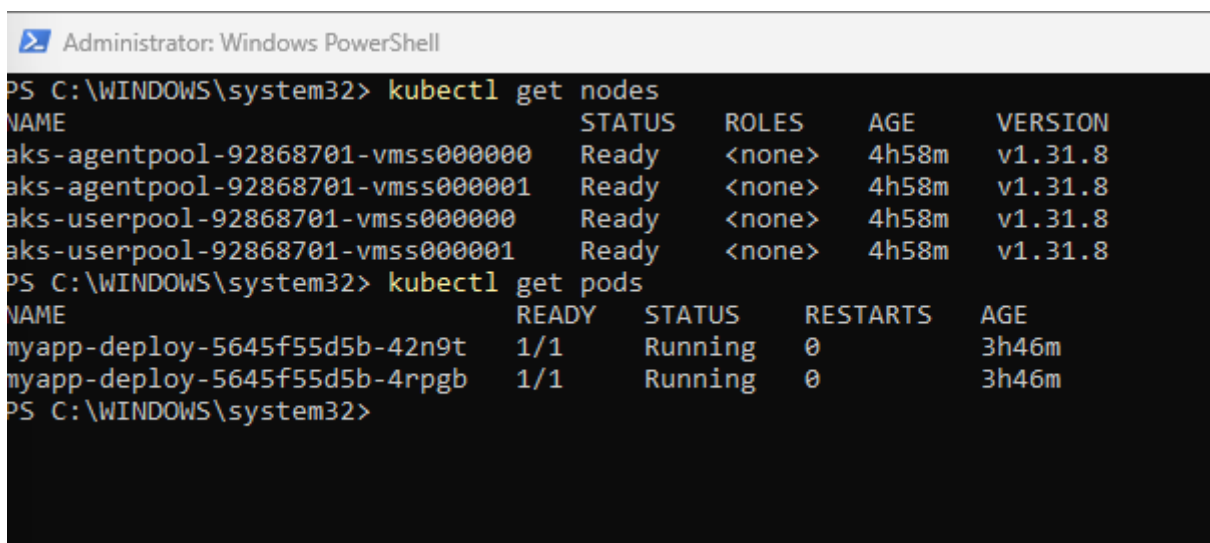
As part of exploring Kubernetes service types, I have already deployed a sample application using the **LoadBalancer service type** on my Azure Kubernetes Service (AKS) cluster.

## REFERENCE TO QUES 1

This setup allowed me to expose the application externally through a **public IP address** assigned by the cloud provider (Azure). This proves the practical use of the LoadBalancer type in real-world, cloud-based environments.

## My Cluster and Pod Setup:

**COMMAND USED :** `kubectl get nodes`  
`kubectl get pods`



```
Administrator: Windows PowerShell
PS C:\WINDOWS\system32> kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
aks-agentpool-92868701-vmss000000  Ready    <none>   4h58m  v1.31.8
aks-agentpool-92868701-vmss000001  Ready    <none>   4h58m  v1.31.8
aks-userpool-92868701-vmss000000    Ready    <none>   4h58m  v1.31.8
aks-userpool-92868701-vmss000001    Ready    <none>   4h58m  v1.31.8
PS C:\WINDOWS\system32> kubectl get pods
NAME                                READY    STATUS    RESTARTS  AGE
myapp-deploy-5645f55d5b-42n9t      1/1      Running   0          3h46m
myapp-deploy-5645f55d5b-4rpgb      1/1      Running   0          3h46m
PS C:\WINDOWS\system32>
```

## YAML File Used for LoadBalancer Service:

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  selector:
```

```
app: myapp
ports:
  - protocol: TCP
    port: 80
    targetPort: 80
type: LoadBalancer
```

## Result:

After applying this YAML, Kubernetes automatically provisioned a public IP through Azure, enabling users to access the application from outside the cluster.

**COMMAND USED :** `kubectl get svc`

This demonstrates how **LoadBalancer services simplify external access** in cloud environments, without needing manual NodePort configuration.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.0.0.1	<none>	443/TCP	5h8m
myapp-service	LoadBalancer	10.0.121.175	4.213.203.76	80:32143/TCP	4h32m

PS C:\WINDOWS\system32>

What the Entries Mean:

kubernetes (ClusterIP):

- This is **automatically created by Kubernetes**.
- It exposes the **Kubernetes API server** internally to the cluster.
- It's **not related to the app**.
- Yes, it is a ClusterIP type service, but **not your own custom one**.

myapp-service (LoadBalancer):

- This is the one **created** for app.
- It has both:
  - A ClusterIP (10.0.121.175) → for **internal access**
  - An External IP (4.213.203.76) → for **public access via LoadBalancer**
- It also auto-assigned a NodePort: 32143 (which is how LoadBalancer works internally).

## What is it?

The `LoadBalancer` type in Kubernetes is used to **expose your service to the internet** by automatically provisioning an **external load balancer** (via the cloud provider like Azure, AWS, or GCP).

## How It Works (Internally):


When you create a service of type `LoadBalancer`:

1. Kubernetes **creates a NodePort** service in the background.
2. The cloud provider provisions a **load balancer with a public IP**.
3. The **load balancer forwards traffic** from its public IP → to the **NodePort on cluster nodes** → which routes to your **Pod(s)**.

## Flow of Traffic:

1. User accesses the **public IP** of the `LoadBalancer`.
2. `LoadBalancer` sends the traffic to one of the **NodePorts**.
3. `NodePort` forwards the request to the **target port** on the Pod.

## Diagram: LoadBalancer Service Flow



## 32143 → Auto-assigned NodePort

- **80** → Target port inside each pod
- Traffic from internet → `LoadBalancer` → `NodePort` → `Pod`

## Advantages of LoadBalancer:

- Simplifies exposure to the internet.
- Automatic provisioning of public IP and traffic routing.
- Ideal for production deployments on cloud platforms like AKS, EKS, GKE

## SO LETS CONTINUE ON THIS AND CREATE A CLUSTER IP SERVICE AND NODEPORT SERVICE FOR THE SAME IMAGE DEPLOYMENT

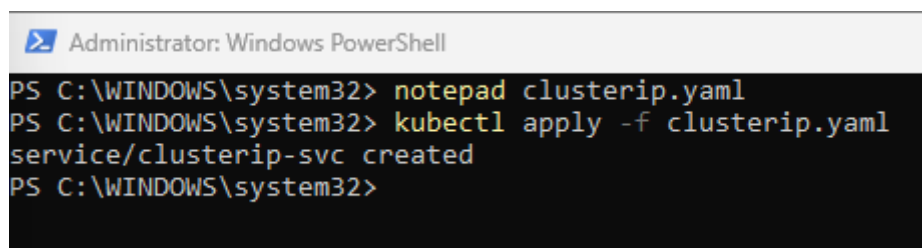
### Step 1: Create and Apply ClusterIP and NodePort Services

#### A. Create `clusterip.yaml`

```
apiVersion: v1
kind: Service
metadata:
  name: clusterip-svc
  labels:
    env: dev
    type: internal
spec:
  selector:
    app: myapp
  ports:
    - port: 80
      targetPort: 80
  type: ClusterIP
```

Apply it:

```
kubectl apply -f clusterip.yaml
```



```
Administrator: Windows PowerShell
PS C:\WINDOWS\system32> notepad clusterip.yaml
PS C:\WINDOWS\system32> kubectl apply -f clusterip.yaml
service/clusterip-svc created
PS C:\WINDOWS\system32>
```

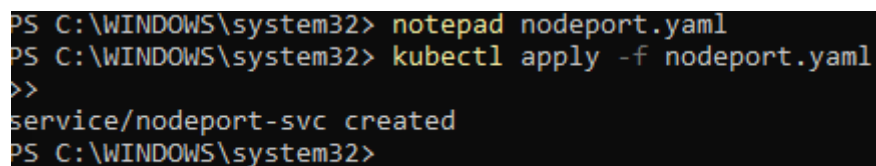
### 3. Create nodeport .yaml

notepad nodeport.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-svc
spec:
  selector:
    app: myapp
  ports:
    - port: 80
      targetPort: 80
      nodePort: 32000
  type: NodePort
```

command used :

kubectl apply -f nodeport.yaml



```
PS C:\WINDOWS\system32> notepad nodeport.yaml
PS C:\WINDOWS\system32> kubectl apply -f nodeport.yaml
>>
service/nodeport-svc created
PS C:\WINDOWS\system32>
```

### 4. Verify Services

kubectl get svc

You should now see 3 services:

- myapp-service (LoadBalancer)
- clusterip-svc (ClusterIP)
- nodeport-svc (NodePort)

```
Administrator: Windows PowerShell
PS C:\WINDOWS\system32> kubectl get svc
>>
NAME                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
clusterip-svc       ClusterIP           10.0.51.13      <none>           80/TCP           3m4s
kubernetes           ClusterIP           10.0.0.1        <none>           443/TCP          5h22m
myapp-service        LoadBalancer        10.0.121.175    4.213.203.76     80:32143/TCP     4h46m
nodeport-svc         NodePort            10.0.100.208    <none>           80:32000/TCP     69s
PS C:\WINDOWS\system32>
```

## 5. Optional: Describe to Check if Pod Is Reached

```
PS C:\WINDOWS\system32> kubectl describe svc clusterip-svc
Name:                clusterip-svc
Namespace:           default
Labels:              <none>
Annotations:         <none>
Selector:            app=myapp
Type:                ClusterIP
IP Family Policy:    SingleStack
IP Families:         IPv4
IP:                 10.0.51.13
IPs:                10.0.51.13
Port:               <unset> 80/TCP
TargetPort:         80/TCP
Endpoints:          10.244.1.241:80,10.244.2.142:80
Session Affinity:    None
Internal Traffic Policy: Cluster
Events:             <none>
PS C:\WINDOWS\system32>
```

```
PS C:\WINDOWS\system32> kubectl describe svc nodeport-svc
Name:                nodeport-svc
Namespace:           default
Labels:              <none>
Annotations:         <none>
Selector:            app=myapp
Type:                NodePort
IP Family Policy:    SingleStack
IP Families:         IPv4
IP:                 10.0.100.208
IPs:                10.0.100.208
Port:               <unset> 80/TCP
TargetPort:         80/TCP
NodePort:           <unset> 32000/TCP
Endpoints:          10.244.1.241:80,10.244.2.142:80
Session Affinity:    None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:             <none>
PS C:\WINDOWS\system32>
```

so until this :

Service	Type	Internal Access	External Access	Status
clusterip-svc	ClusterIP	✓ Yes (within cluster)	✗ No	Working
nodeport-svc	NodePort	✓ Yes (if port-forwarded)	✗ No (no external node IP)	Working
myapp-service	LoadBalancer	✓ Yes	✓ Yes (4.213.203.76)	Working

WE CANNOT ACCESS PUBLICALLY CAUSE :

NodePort mapped the app to a static port (32000) on each node, but due to lack of external node IPs in AKS, it cannot be accessed directly from the internet.

### WHAT WE CAN DO TO ACCESS IT PUBLICALLY:

**Option 1:** Manually Expose a VM Node to the Internet (NOT recommended)

Steps (Advanced & Not Secure for Production):

1. Create a **public IP address** in Azure.
2. Assign that IP to a **network interface** on a VM Scale Set (one of the AKS nodes).
3. Open firewall port (e.g., 32000).
4. Then you can access:

`http://<Public-Node-IP>:<NodePort>`

### Option 2: Use a LoadBalancer to Simulate NodePort (Cleanest + Practical)

Since Azure doesn't expose nodes directly, you can:

1. Keep your service as NodePort
2. **Manually create a LoadBalancer** resource in Azure Portal



3.Point it to backend pool → Node IPs → NodePort

4.Result: Public IP hits NodePort behind the scenes

**But this defeats the purpose because...**

Kubernetes LoadBalancer service already does this automatically!

SO ITS DONE FINALLY:

In this task, I explored all three main types of Kubernetes Services: ClusterIP, NodePort, and LoadBalancer.

I deployed a sample app (myapp-deploy) and connected it to all 3 services:

- ClusterIP was used to enable internal communication inside the cluster.
- NodePort mapped the app to a static port (32000) on each node, but due to lack of external node IPs in AKS, it cannot be accessed directly from the internet.
- LoadBalancer was used to expose the app to the public internet via Azure's managed load balancer, which provided a working external IP (4.213.203.76).

SO WE CAN UNDER STAND THE USE CASES AS :

## Real-World Use Case Comparison of Kubernetes Service Types

### ◆ 1. ClusterIP - Internal Service Communication Only

#### Theory:

ClusterIP exposes the service **only inside the Kubernetes cluster**, typically used for **backend-to-backend** communication between microservices.

#### Real-World Example:

In a real e-commerce site, the **frontend service** (React or Angular) calls the **backend service** (Node.js or Django) via ClusterIP. The backend in turn calls **database** or **authentication services**, all of which are invisible outside the cluster.

#### Why?

These components don't need to be public. Using ClusterIP keeps them secure and isolated.

### 2. NodePort - Debug or Dev-Level External Access

#### Theory:

NodePort exposes the service on a **fixed port (30000-32767)** on every node's IP. It's often used for **quick testing or temporary external access**, especially in **bare metal or local setups**.

#### Real-World Example:

In a private company network, a dev team may use NodePort to expose a service to their **office network** so team members can test the staging build by visiting `http://<node-ip>:32000`.

#### Why?

No need for a full cloud LoadBalancer in non-production or internal settings. But not ideal for public or large-scale access.

### 3. LoadBalancer – Production-Grade Public Access

#### Theory:

LoadBalancer provisions a **cloud-managed external IP** and routes public traffic to the app. It's best for production websites and APIs exposed to the internet.

#### Real-World Example:

In real production apps like **Tabcura, Amazon, or Netflix**, frontend services are exposed via LoadBalancer.

For instance:

<https://www.netflix.com> → hits a **cloud LoadBalancer** → routes to Kubernetes pods hosting their frontend.

#### Why?

It provides scalability, public reach, and is managed by the cloud (e.g., Azure Load Balancer in AKS).