

QUES 3 : PersistentVolume (PV) and PersistentVolumeClaim (PVC)

SOLN:

In Kubernetes, a **PersistentVolume (PV)** is a piece of storage in the cluster that has been provisioned by an admin or dynamically provisioned using StorageClasses. A **PersistentVolumeClaim (PVC)** is a request for that storage by a user. This mechanism allows containers to store data persistently, even across pod restarts or rescheduling.

In simple words, **PersistentVolume (PV)** is like a hard disk that's kept aside for use in your Kubernetes system. And **PersistentVolumeClaim (PVC)** is like a request your app makes saying, "Hey, I need some space to store my data."

Normally, when a pod stops or restarts, it loses all its data. But when you use PV and PVC, your app gets permanent storage — like saving your files on a pen drive instead of RAM, so they don't get lost.

In the previous tasks (Q1 & Q2), I successfully deployed a sample application (myapp-deploy) and exposed it via a Kubernetes service. Now, in this task (Q3), I am implementing **persistent storage** using **PersistentVolume (PV)** and **PersistentVolumeClaim (PVC)** to ensure the application can store and retain data even if the pod restarts or gets rescheduled.

```
PS C:\WINDOWS\system32> kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
aks-agentpool-92868701-vmss000000  Ready    <none>   6h16m  v1.31.8
aks-agentpool-92868701-vmss000001  Ready    <none>   6h16m  v1.31.8
aks-userpool-92868701-vmss000000    Ready    <none>   6h16m  v1.31.8
aks-userpool-92868701-vmss000001    Ready    <none>   6h16m  v1.31.8
PS C:\WINDOWS\system32> kubectl get pod
NAME                                READY    STATUS    RESTARTS  AGE
myapp-deploy-5645f55d5b-42n9t      1/1      Running   0          5h3m
myapp-deploy-5645f55d5b-4rpgb      1/1      Running   0          5h3m
PS C:\WINDOWS\system32> kubectl get pods
NAME                                READY    STATUS    RESTARTS  AGE
myapp-deploy-5645f55d5b-42n9t      1/1      Running   0          5h3m
myapp-deploy-5645f55d5b-4rpgb      1/1      Running   0          5h3m
PS C:\WINDOWS\system32>
```

ALL THE DEPLOYMENTS ARE WORKING FINE

LETS CREATE A YAML FILE TO CONFIGURE PV & PVC FOR THE PROJECT:

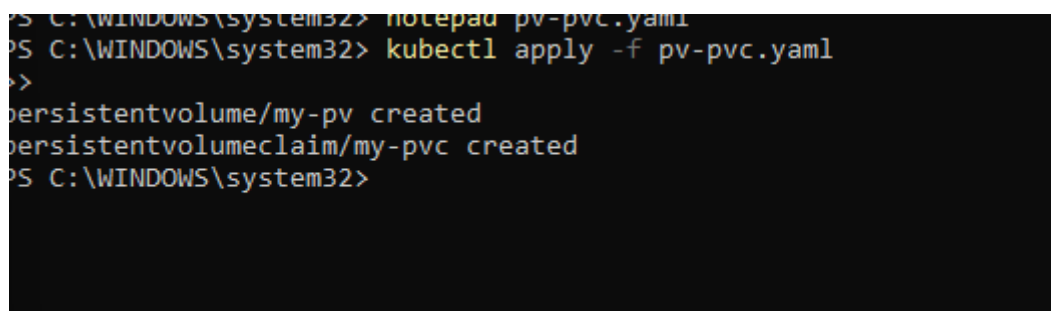
pv-pvc.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce
  storageClassName: manual
  hostPath:
    path: "/mnt/data" # For demo/testing only
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: manual
  resources:
    requests:
      storage: 1Gi
```

lets apply it :

```
kubectl apply -f pv-pvc.yaml
```




```
PS C:\WINDOWS\system32> notepad pv-pvc.yaml
PS C:\WINDOWS\system32> kubectl apply -f pv-pvc.yaml
>
persistentvolume/my-pv created
persistentvolumeclaim/my-pvc created
PS C:\WINDOWS\system32>
```

PVC AND PV HAS BEEN SUCCESSFULLY CREATED

LETS TRY TO MOUNT THESE PVC AND PV TO OUR PREVIOUS BUILT DEPLOYMENT

STEP : MODIFY THE DEPLOYMENT.YAML FILE

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deploy
spec:
  replicas: 2
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      containers:
        - name: myapp-container
          image: prateek2004/my-frontend
          ports:
            - containerPort: 80
          volumeMounts:
            - name: app-storage
              mountPath: /app/data #  You can change this to wherever your app writes
files
      volumes:
        - name: app-storage
          persistentVolumeClaim:
            claimName: my-pvc
```

NOTE:

`volumeMounts` tells your container to use storage inside the container at `/app/data`.

- `volumes` connects the PVC (my - pvc) to the pod.
- You can change `/app/data` to wherever your frontend app writes data (like user uploads, logs, etc.).

Save and close.

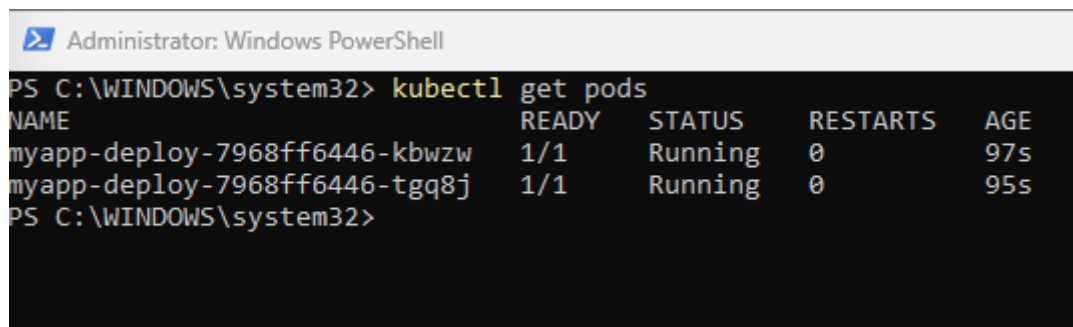
3. Then reapply it:

`kubectl apply -f deployment.yaml`

```
error: the path "myapp-deploy.yaml" does not exist
PS C:\WINDOWS\system32> kubectl apply -f deployment.yaml
deployment.apps/myapp-deploy configured
PS C:\WINDOWS\system32>
```

Verify It Worked:

`kubectl get pods`

A screenshot of a Windows PowerShell terminal window titled "Administrator: Windows PowerShell". The terminal shows the command `kubectl get pods` being executed. The output is a table with five columns: NAME, READY, STATUS, RESTARTS, and AGE. There are two rows of pod information.

NAME	READY	STATUS	RESTARTS	AGE
myapp-deploy-7968ff6446-kbwzw	1/1	Running	0	97s
myapp-deploy-7968ff6446-tgq8j	1/1	Running	0	95s

The terminal prompt is `PS C:\WINDOWS\system32>`.

`kubectl describe pod <your-pod-name>`

```

C:\WINDOWS\system32> kubectl describe pod myapp-deploy-7968ff6446-kbwzw
Name: myapp-deploy-7968ff6446-kbwzw
Namespace: default
Priority: 0
Service Account: default
Node: aks-userpool-92868701-vmss000001/10.224.0.4
Start Time: Mon, 30 Jun 2025 20:13:03 +0530
Labels: app=myapp
        pod-template-hash=7968ff6446
Annotations: <none>
Status: Running
IP: 10.244.1.214
Controlled By: ReplicaSet/myapp-deploy-7968ff6446
Containers:
  myapp-container:
    Container ID: containerd://b308878b8b7b0940475d3a22e41212810fc78451338e0d3d012c4cbd1ff75508
    Image: prateek2004/my-frontend
    Image ID: docker.io/prateek2004/my-frontend@sha256:58bedad0762aca5ffa1d2e7f2f9d275b5677bca263ceb0deed5cca675888b7a
    Port: 80/TCP
    Host Port: 0/TCP
    State: Running
      Started: Mon, 30 Jun 2025 20:13:05 +0530
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /app/data from app-storage (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-x447p (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized         True
  Ready               True
  ContainersReady     True
  PodScheduled        True
Volumes:
  app-storage:
    Type: PersistentVolumeClaim (a reference to a PersistentVolumeClaim in the same namespace)
    ClaimName: my-pvc
    ReadOnly: false
  kube-api-access-x447p:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    Optional: false
    DownwardAPI: true
OS Class: BestEffort
Node-Selectors: <none>
Operations:
  node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
  node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   2m18s  default-scheduler  Successfully assigned default/myapp-deploy-7968ff6446-kbwzw to aks-use

```

• we can see a successfull mount

kubectl exec -it <your-pod-name> -- sh

ls /app/data

echo "Testing PVC connection" > /app/data/test.txt

cat /app/data/test.txt

HENCE ITS SUCCESSFULLY VERIFIED AND TESTED:

```
Administrator: Windows PowerShell
PS C:\WINDOWS\system32> kubectl exec -it myapp-deploy-7968ff6446-kbwzw -- sh
/ # ls /app/data
/ # pwd
/
/ # ls /app/data
/ # echo "testing connection" > /app/data/test.txt
/ # cat /app/data/test.txt
testing connection
/ #
```

In real-world applications, PersistentVolumes and PersistentVolumeClaims are used wherever long-term storage is required — such as databases (MySQL, MongoDB), file upload systems, media platforms, and content management systems like WordPress. Without persistent storage, all application data would be lost every time a pod is recreated or updated.